

Operator scaling++

Yuanzhi Li (Princeton)

Joint work with Zeyuan Allen-Zhu, Ankit Garg(Microsoft Research),
Rafael Oliveria (University of Toronto) and Avi Wigderson (IAS)

Thanks:

My fiancée Yandi Jin

Problem of interest

- Group action: A group G , a set V .
 - Each element g in G defines a map $g: V \rightarrow V$ (actions).
 - The actions **preserve** the group structure:
 - $(g_1 g_2)(v) = g_1(g_2(v))$.
 - $\text{Identity}(v) = v$.

Problem of interest

- Orbit intersection
 - Given elements v_1, v_2 in V , whether there are g_1, g_2 in G such that
 - $g_1(v_1) = g_2(v_2)$.
- Orbit closure intersection:
 - Suppose V is some metric space.
 - Given elements v_1, v_2 in V , consider two sets:
 - $O(v_1) = \{ g(v_1) \mid g \text{ in } G \}$, $O(v_2) = \{ g(v_2) \mid g \text{ in } G \}$.
 - Whether the **closures** of $O(v_1)$ and $O(v_2)$ intersect.
- We want to do it efficiently.

Why is this problem interesting?

- Graph isomorphism:
 - G : Symmetric group on n elements.
 - V : Set of graphs with n vertices.
- An efficient algorithm (poly time) would be a huge break through.

How hard is this problem?

- Orbit intersection: Unlikely to be NP-complete.
- More like PIT (via invariant polynomials).

In this paper

- We give a (deterministic) polynomial time algorithm for orbit closure intersection problem for left-right linear actions:
 - V = set of m tuples of $n \times n$ complex matrices.
 - $v = (M_1, \dots, M_m)$.
 - $G = SL_n(\mathbb{C}) \times SL_n(\mathbb{C})$.
 - $SL_n(\mathbb{C})$: The set of all linear operations on \mathbb{C}^n with determinate 1.
 - $g = A, B$.
 - $g(v) = (AM_1B^+, \dots, AM_mB^+)$.
- Invariant polynomial: $p(X_1, \dots, X_m) = \det(X_1 \otimes M_1 + \dots + X_m \otimes M_m)$

The main problem:

- Given $v_1 = (M_1, \dots, M_m)$, $v_2 = (M_1', \dots, M_m')$, decide whether for every $\epsilon > 0$ there exists A, B in such that
 - $\|AM_1B^+ - M_1'\|_F^2 + \dots + \|AM_mB^+ - M_m'\|_F^2 < \epsilon$.
- We give a polynomial time algorithm:
 - Polynomial in the **number of bits** to write down v_1, v_2 .

Optimization approach

- Let us focus on $V = \mathbb{C}^n$ (complex Euclidean space of dimension n), G is a subset of $GL_n(\mathbb{C})$.
- $\| \cdot \|_2$ is the Euclidean norm.
 - Define $f_v(g) = \|g(v)\|_2^2$.
 - Define $N(v) = \inf_{g \in G} f_v(g)$.
 - Moment map: $\mu_G(v)$.

Optimization approach

- Kempf-Ness Theorem + Hilbert Nullstellensatz:
 - Suppose $N(v_1), N(v_2) > 0$, then there is an element in the intersection of the closures of $O(v_1)$ and $O(v_2)$:
 - Then there is one v_0 in the intersection of the closures of $O(v_1)$ and $O(v_2)$ such that:
 - $\mu_G(v_0) = 0$.
 - v_0 is **unique** in the sense that for every such v_0, v_0' , there exists s with:
 - $v_0 = s(v_0')$, where:
 - $||s(v')||_2 = ||v'||_2$ for every v' in V .
 - s is in the maximum compact subgroup K of G .

Two steps approach

- Find $v_1' = g_1(v_1)$, $v_2' = g_2(v_2)$ such that $\mu_G(v_1') = \mu_G(v_2') = 0$.
 - Find g_1, g_2 that minimizes $\|g_1(v_1)\|_2^2, \|g_2(v_2)\|_2^2$.
 - Optimization: given v , find the **argmin** of $\|g(v)\|_2^2$ over g in G .
- Solve the problem whether the orbit of v_1', v_2' in **K** intersects
 - Solve the problem on a simpler group (The maximum compact subgroup).
- Original \leftrightarrow Opt + Simple.

The idea looks simple

- But there is a problem:
 - The optimization step: Given v , find the argmin of $\|g(v)\|_2^2$ over g in G .
 - Usually, the theorems in optimization looks like this:
 - Given v , find g' such that $\|g'(v)\|_2^2 \leq \inf_{g \in G} \|g(v)\|_2^2 + \epsilon$.
 - In some $\text{Time}(1/\epsilon)$.
- Two reasons:
 - Infimum might not be achievable.
 - Infimum might not be a rational matrix (even when v is an integer vector).

Can we work inexactly?

- One major difference between optimization and mathematics:
 - In math:
 - The exact minimizer has property blah blah blah.
 - If \$ equals to ¥, then blah blah blah.
 - In optimization:
 - To get an efficient algorithm, most of the time we need to work with:
 - The approximate minimizer.
 - When \$ approximately equals to ¥.

Two steps approach (Modified)

- Find g_1, g_2 that approximately minimizes $||g_1(v_1)||_2^2, ||g_2(v_2)||_2^2$
 - Optimization: given v , **approximately** minimizes $||g(v)||_2^2$ over g in G .
 - $||g'(v)||_2^2 \leq \inf_{g \in G} ||g(v)||_2^2 + \epsilon.$
- Solve the problem whether the orbit of v_1', v_2' in K **approximately** intersects:
 - Find s_1, s_2 in K such that:
 - $||s_1(v_1') - s_2(v_2')||_2^2 \leq \epsilon.$
- Original \leftrightarrow Opt(ϵ) + Simple (ϵ)?

How to choose epsilon?

- How fast we can minimize the given function?
 - Given v , find g' such that $\|g'(v)\|_2^2 \leq \inf_{g \in G} \|g(v)\|_2^2 + \epsilon$
 - In some $\text{Time}(1/\epsilon)$.
 - What is this $\text{Time}(1/\epsilon)$?
 - Logarithmic? $\log(1/\epsilon)$?
 - Polynomial? $(1/\epsilon)^{10}$?
 - Exponential? $e^{1/\epsilon}$?
 - \rightarrow we want large epsilon so we can optimize fast.
- How small the error needs so the mathematical theorems still hold:
 - \rightarrow we want small epsilon so we can prove the theorem easily.

How fast we can minimize the given function

- Given v , minimize $\|g(v)\|_2^2$ over g in G .
- Theorem of [GRS'13, Wood'11]
 - For 'most of the' G : a subset of $GL_n(\mathbb{C})$, $V = \mathbb{C}^n$:
 - $\|g(v)\|_2^2$ is geodesically convex.

Geodesic Convexity

- Equip some Riemannian metric $|| \cdot ||^*$ on G : a subset of $GL_n(\mathbb{C})$.
 - Geodesic path from g_1 to g_2 :
 - $\gamma: [0, 1] \rightarrow G$
 - $\gamma(0) = g_1$
 - $\gamma(1) = g_2$
 - For every s, t in $[0, 1]$: $|| \gamma(s) - \gamma(t) ||$ is **proportional** to $|s - t|$.
 - It can also be (uniquely) characterized by **$\gamma'(0)$** .
- A function $f: G \rightarrow \mathbb{R}$ is geodesically convex:
 - If and only if for every geodesic path γ :
 - $f(\gamma): [0, 1] \rightarrow \mathbb{R}$ is convex.

Minimize a geodesically convex function

- Convex function f :
 - Gradient descent: at each point x , move along $-\nabla f(x)$.
- Geodesic convex function f :
 - Geodesic gradient descent: At each point x , find geodesic path γ such that
 - $\gamma(0) = x$.
 - $\gamma'(0) = -\nabla f(x)$.
 - Move along γ .

Minimize a Geodesically convex function

- Theorem[ZS'16]:
 - Using Geodesic gradient descent, we can minimize a geodesic convex function f up to error ϵ in time:
 - $(1/\epsilon)^2$.
 - Polynomial time algorithm for (inverse) polynomially small ϵ .

Mathematical part

- Can we prove the theorem when the error ϵ is inverse polynomial?
 - It is ok for applications in previous work: [GGOW'16] (Null Cone)
 - **Not ok** for our problem: We need exponentially small epsilon.

Mathematical part

- The inexact orbit closure intersection problem:
 - Given tuples $v_1 = (M_1, \dots, M_m)$ and $v_2 = (M_1', \dots, M_m')$.
 - Our group $G: SL_n(\mathbb{C}) \times SL_n(\mathbb{C})$.
 - The orbit closure of v_1 and v_2 intersects:
 - if and only if there exists (A, B) in G such that
 - $||AM_1B^+ - M_1'||_F^2 + \dots + ||AM_mB^+ - M_m'||_F^2 \leq \epsilon$.
- Central question: How small ϵ needs to be?

Mathematical part

- Theorem (this paper): epsilon being (inverse) exponentially small is sufficient:
- There exists a **polynomial** $p: \mathbb{R}^3 \rightarrow \mathbb{R}$ such that for every m, n, B , every $v_1 = (M_1, \dots, M_m)$ and $v_2 = (M_1', \dots, M_m')$ where M_i, M_i' in $\mathbb{C}^{n \times n}$ with each entry being integers in $[-B, B]$, the following two statements are equivalent:
 - The orbit closure of v_1 and v_2 intersects.
 - There exists (A, B) in G such that $\|AM_1B^+ - M_1'\|_F^2 + \dots + \|AM_mB^+ - M_m'\|_F^2 \leq e^{-p(m, n, \log B)}$.
- (inverse) exponential is tight.

Epsilon is inverse exponentially small

- To get an efficient algorithm, we need to:
 - Given v , find g' such that $\|g'(v)\|_2^2 \leq \inf_{g \in G} \|g(v)\|_2^2 + \epsilon$.
 - In time $\text{polylog}(1/\epsilon)$.
- Which means that we **can not** use gradient descent.

Faster optimization algorithms

- Optimization algorithm with $\text{polylog}(1/\epsilon)$ convergence rate:
 - In convex setting:
 - Newton's method (Only local convergence rate).
 - Interior point algorithm.
 - Ellipsoid algorithm.

Faster optimization algorithms:

- Newton's method (Only local convergence rate).
- Interior point algorithm.
 - They are 'gradient descent' type of algorithms.
 - They all based on the so called ``self-concordant function''.

Self-concordant function:

- A convex function $f: \mathbb{R} \rightarrow \mathbb{R}$ is self-concordant if for every x :
 - $|f'''(x)| \leq 2 (f''(x))^{3/2}$.
 - **Scaling independent:** $|f'''(ax)| \leq 2 (f''(ax))^{3/2}$ for every non-zero a .
- A geodesically convex function is self-concordant if for every geodesic path γ , $f(\gamma): \mathbb{R} \rightarrow \mathbb{R}$ is self-concordant.

Self-concordant function:

- A convex function $f: \mathbb{R} \rightarrow \mathbb{R}$ is self-concordant if for every x :
 - $|f'''(x)| \leq 2 (f''(x))^{3/2}$
- 'looks like' Quadratic function:
 - When the second order derivative is small, the change of it is also small.

Optimize a self-concordant function:

- Forklore (informal): There is an algorithm that runs in time $\text{poly}(n)\log(1/\epsilon)$ to minimize a geodesic self-concordant function: $G \rightarrow \mathbb{R}$ up to accuracy ϵ .
- Move along direction $-(\nabla^2 f(x))^{-1} \nabla f(x)$.

Good, can we use it?

- No...
 - Our function $\|g(v)\|_2^2$ is **not** geodesic self-concordant ...

Ok, can we modify the definition?

- $|f'''(x)| \leq 2 (f''(x))^{3/2}$
- Why **3/2**? I don't like 3/2... it's not an integer... not nice...
- Scaling independent.
- Can we fix a scaling and use other powers?

Self-robust function:

- A function $f: \mathbb{R} \rightarrow \mathbb{R}$ is self-robust if for every x :
 - $|f'''(x)| \leq f''(x)$.
- A geodesic convex function is self-concordant if for every unit speed geodesic path γ , $f(\gamma): \mathbb{R} \rightarrow \mathbb{R}$ is self-concordant.
 - Unit speed: $\|\gamma'(0)\| = 1$.

Optimize a self-robust function:

- Theorem [This paper, informal]: There is an algorithm that runs in time $\text{poly}(n)\log(1/\epsilon)$ to minimize a geodesic self-robust function: $G \rightarrow R$ up to accuracy ϵ . (G is a subset of $GL_n(\mathbb{C})$)

Overview of the algorithm:

- At every iteration, maintain a point g in G .
 - Compute the local geodesic gradient $\nabla f(g)$, defined as:
 - For every direction e , $\langle e, \nabla f(g) \rangle = f'(\gamma(t))|_{t=0}$
 - Such that $\gamma(0) = g$, $\gamma'(0) = e$.
 - Compute the local geodesic hessian $\nabla^2 f(g)$, defined as
 - For every direction e , $e^+ \nabla^2 f(g) e = f''(\gamma(t))|_{t=0}$
 - Such that $\gamma(0) = g$, $\gamma'(0) = e$.
 - Minimize the function $g(e) = \langle e, \nabla f(g) \rangle + 0.1 e^+ \nabla^2 f(g) e$ over $\|e\| \leq 0.1$
 - Let e^* be the minimizer.
 - Move to $\gamma(0.01)$ such that
 - $\gamma(0) = g$, $\gamma'(0) = e^*$.

Wait a second...

- Why $\nabla f(g)$, $\nabla^2 f(g)$ even exist?
- It is important that γ is geodesic path.
 - Can be defined via 'Exponential map'.

Good, can we use it?

- No...
 - Our function $\|g(v)\|_2^2$ is not geodesic self-robust...

Wait... so what are you talking about???

Modify the function

- We just need to find the minimizer g in G of $\|g(v)\|_2^2$.
- We can minimize any function $h_v(g)$ such that
 - $\operatorname{argmin}_{g \in G} h_v(g) = \operatorname{argmin}_{g \in G} \|g(v)\|_2^2$.
- Can we find such function? And make it self-concordant/self-robust?

The log capacity function:

- In our problem:
 - $v = (M_1, \dots, M_m)$, $g = (A, B)$.
 - $\|g(v)\|_2^2 = \|AM_1B^+\|_F^2 + \dots + \|AM_mB^+\|_F^2$.
 - Minimizing this function is aka operator scaling.
- The equivalent function: **log capacity function**: for a PSD matrix X ,
 - $f(X) = \log \det(M_1XM_1^+ + \dots + M_mXM_m^+) - \log \det(X)$.
- Theorem [Gurvits'04]:
 - Let X^* be a minimizer of $f(X)$, let A^* , B^* be the minimizer of $\|g(v)\|_2^2$, then there exists a, b in \mathbb{R} such that
 - $B^*B^* = a X^*$.
 - $A^*A^* = b(M_1XM_1^+ + \dots + M_mXM_m^+)$.

Log capacity function:

- Theorem [This paper]:
- $f(X) = \log\det(M_1XM_1^+ + \dots + M_mXM_m^+) - \log\det(X)$
 - Is a geodesic **self-robust** function, with the geodesic path over PSD matrices given as:
 - $\gamma(t) = X_0^{1/2}e^{tD}X_0^{1/2}$
 - $\gamma(0) = X_0$
 - $\gamma'(0) = X_0^{1/2}DX_0^{1/2}$.

Step 1:

- Minimize:
 - $f_1(X) = \log\det(M_1 X M_1^+ + \dots + M_m X M_m^+) - \log\det(X)$.
 - $f_2(X) = \log\det(M_1' X M_1'^+ + \dots + M_m' X M_m'^+) - \log\det(X)$.
 - Let X_1^*, X_2^* be an ϵ -approximate minimizers.
- Let:
 - $B^+ B = a_1 X_1^*, B_2^+ B_2 = a_2 X_2^*$.
 - $AA^+ = b_1(M_1 X_1^* M_1^+ + \dots + M_m X_1^* M_m^+), A_2 A_2^+ = b_2(M_1' X_2^* M_1'^+ + \dots + M_m' X_2^* M_m'^+)$.

Step 2:

- Now, let $w_1 = Av_1B^+$, $w_2 = A_2v_2B_2^+$.
- Check the orbit closure w_1, w_2 approximately interests in a subgroup K :
 - K : all the elements g in G such that for every v in V , $\|g(v)\|_2^2 = \|v\|_2^2$
 - In this problem: K : the set of all (determinate one) **unitary** matrices.

Exact unitary equivalence:

- Given $w_1 = (M_1, \dots, M_m)$, $w_2 = (M_1', \dots, M_m')$,
- Check whether there exists unitary matrices U, V such that
 - For all i in m : $U M_i V^+ = M_i'$.
- Existing algorithms [CIK'97, IQ'18] can solve it in polynomial time.

Inexact unitary equivalence:

- Given $w_1 = (M_1, \dots, M_m)$, $w_2 = (M_1', \dots, M_m')$,
- Check whether there exists unitary matrices U, V such that
 - For all i in m : $\|U M_i V^+ - M_i'\|_F \leq \epsilon$.
 - Where ϵ is (inverse) exponentially small.

Naïve idea

- Given $w_1 = (M_1, \dots, M_m)$, $w_2 = (M_1', \dots, M_m')$,
- Check whether there exists unitary matrices U, V such that
 - For all i in m : $\|U M_i V^\dagger - M_i'\|_F \leq \epsilon$.
 - Where ϵ is (inverse) exponentially small.
 - Make ϵ **smaller** than (inverse) exponential of the bit complexity of M_i and M_i' .
 - So $U M_i V^\dagger = M_i'$.
 - Use exact algorithm.

Recall how we get w_1

- Minimize:

- $f_1(X) = \log\det(M_1 X M_1^+ + \dots + M_m X M_m^+) - \log\det(X)$.
- $f_2(X) = \log\det(M_1' X M_1'^+ + \dots + M_m' X M_m'^+) - \log\det(X)$.
- Let X_1^* , X_2^* be an ϵ -approximate minimizers.

- Let:

- $B^+ B = a_1 X_1^*$, $B_2^+ B_2 = a_2 X_2^*$.
- $AA^+ = b_1(M_1 X_1^* M_1^+ + \dots + M_m X_1^* M_m^+)$, $A_2 A_2^+ = b_2(M_1' X_2^* M_1'^+ + \dots + M_m' X_2^* M_m'^+)$.

Naïve idea

- Given $w_1 = (M_1, \dots, M_m)$, $w_2 = (M_1', \dots, M_m')$:
 - $w_1 = Av_1B^+$, where A, B defined by an ϵ -approximate minimizer.
- The smaller ϵ is, the larger the bit complexity of M_i is.
 - In fact, ϵ can **never** be smaller than (inverse) exponential of the bit complexity of M_i and M_i' .

This paper

- We give an algorithm that runs in time $\text{poly}(m, n, \log 1/\epsilon, B)$ that
 - Given $w_1 = (M_1, \dots, M_m)$, $w_2 = (M_1', \dots, M_m')$,
 - Distinguish whether there exists unitary matrices U, V such that
 - For all i in m : $\|U M_i V^\dagger - M_i'\|_F \leq \epsilon$.
 - There exists i in m : $\|U M_i V^\dagger - M_i'\|_F \geq \exp(\text{poly}(m, n))\epsilon^{1/\text{poly}(m, n)}$.
- It applies to **any** ϵ , regardless of the bit complexity B of w_1, w_2 .

In exact algorithm: Step 1

- What do we know if there exists unitary U, V such that
 - $\|U M V^+ - M'\|_F \leq \epsilon$.
- Singular value decomposition.

Singular value decomposition (SVD)

- For every matrix M in $\mathbb{C}^{n \times n}$:
- There exists unitary matrices U , V and a diagonal matrix $\Sigma = \text{diag}(\sigma_1, \dots, \sigma_n)$ with $\sigma_1 \geq \sigma_2 \geq \dots \geq \sigma_n \geq 0$ such that:
 - $M = U\Sigma V^+$.

Gap-free Wedin Theorem

- What do we know if there exists unitary U, V such that
 - $\|U M V^+ - M'\|_F \leq \epsilon$?
- Theorem [AL'16]: Let $M = U_1 \Sigma V_1^+$, $M' = U_2 \Sigma V_2^+$ be the SVD of M, M' . Suppose $\|U M V^+ - M'\|_F \leq \epsilon$. For every $\delta \geq 0$, if
 - There exists k in $[n - 1]$ such that $\sigma_k - \sigma_{k+1} \geq \delta$
 - Then $\|U_{11}^+ U^+ U_{22}\|_F \leq \epsilon/(\delta - 2\epsilon)$.
 - U_{11} : First k columns of U_1 , U_{22} : Last $n - k$ columns of U_2

Gap-free Wedin Theorem

- $\|U_{11}^* U^* U_{22}\|_F \leq \epsilon/(\delta - 2\epsilon).$
- If we write $U = U_2 U' U_1^*$ for an unitary matrix U' .
 - Then U' is close to being (block) diagonal.

Apply gap-free Wedin Theorem

- Suppose $\|U M_i V^+ - M_i'\|_F \leq \epsilon$ and M_i has a singular value gap δ
 - Then we can find in polynomial time unitary U', U'', V', V'' such that there exists unitary matrices U_1, U_2, V_1, V_2 :
 - $\|U - U' \text{diag}(U_1, U_2) U''^+\|_F \leq \epsilon/(\delta - 2\epsilon)$.
 - $\|V - V' \text{diag}(V_1, V_2) V''^+\|_F \leq \epsilon/(\delta - 2\epsilon)$.
 - Reduce the original problem to **two sub problems** of smaller dimensions.

In the end:

- Left with a problem where all M_i has no singular value gap:
 - M_i is close to a (constant multiple of) unitary matrix.
- $\|U M_i V^+ - M_i'\|_F \leq \epsilon$ where M_i, M_i' are close to unitary:
 - V is close to $M_i'^+ U M_i$.
- So we can reduce one unitary V and focus **only on U** .

Inexact unitary conjugation:

- Given $w_1 = (M_1, \dots, M_m)$, $w_2 = (M_1', \dots, M_m')$,
- Check whether there exists unitary matrices U such that
 - For all i in m : $\|U M_i U^+ - M_i'\|_F \leq \epsilon$.
 - Where ϵ is (inverse) exponentially small.

Eigenvalue

- The eigenvalues of a matrix M is given by the set of all values λ in \mathbb{C} such that
 - $\text{Det}(\lambda I - M) = 0$.

Eigenvalue Wedin Theorem

- What do we know if there exists unitary U such that
 - $\|U M U^+ - M'\|_F \leq \epsilon$.
- Theorem [This paper]: Suppose $\|U M U^+ - M'\|_F \leq \epsilon$, then for every $\delta \geq 0$, if M has two eigenvalues λ_1, λ_2 such that $|\lambda_1 - \lambda_2| \geq \delta$,
 - We can compute, in time $\text{poly}(n, \log 1/\epsilon)$ unitary matrices U', U'' such that there exists unitary matrices U_1, U_2 :
 - $\|U - U' \text{diag}(U_1, U_2) U''^+ \|_F \leq \epsilon/(\delta/n - 2\epsilon)$.

Apply eigenvalue Wedin Theorem

- Suppose $\|U M_i U^+ - M_i'\|_F \leq \epsilon$ and M_i has a eigenvalue gap δ :
 - Then we can find in polynomial time unitary matrices U', U'' such that there exists unitary matrices U_1, U_2
 - $\|U - U' \text{diag}(U_1, U_2) U''^+ \|_F \leq \epsilon/(\delta/n - 2\epsilon)$.
- Reduce the original problem to **two sub problems** of smaller dimensions.

In the end:

- Left with a problem where all M_i has no singular value gap nor eigenvalue gap.
 - Theorem [This paper]: M_i is close to (constant multiple of) identity.

Summary

- We give a polynomial time algorithm for orbit closure intersection problem for left-right linear actions:
 - V = set of m tuples of $n \times n$ complex matrices.
 - $v = (M_1, \dots, M_m)$.
 - $G = \mathrm{SL}_n(\mathbb{C}) \times \mathrm{SL}_n(\mathbb{C})$.
 - $g = A, B$.
 - $g(v) = (AM_1B^+, \dots, AM_mB^+)$.

Summary

- Use geodesic optimization to reduce the original problem to inexact unitary equivalence problem.
 - Inexact theorem holds for (inverse) exponentially small epsilon.
 - Mathematics.
 - Design an algorithm with linear convergence rate.
 - Optimization.
- Design a new algorithm for inexact unitary equivalence problem.